

IaCSecBench: A Finding-Normalization Methodology and Reproducible Harness for Evaluating Infrastructure-as-Code Security Validation

Manideep Chittineni (ORCID:
0009-0003-9709-5842)

Received: date / Accepted: date

Abstract Context. Infrastructure as Code (IaC) dominates cloud provisioning, yet pre-deployment security validation remains fragmented across static scanners, policy engines, and test frameworks. Incompatible abstraction levels, bespoke schemas, and developer-authored benchmarks impede rigorous comparison. **Objective.** We investigate how far the matching criterion (what counts as a detection) determines a comparison’s conclusions. **Method.** IaCSecBench is an open evaluation harness that normalizes heterogeneous scanner outputs onto a canonical control taxonomy, enforces mechanical corpus admissibility, and scores under three matching criteria. Over 56 admissible cases (30 vulnerable, 26 compliant), we evaluate three third-party scanners, plan-level policy evaluation, and a repository-edge layer, reporting exact Clopper–Pearson intervals and exact McNemar tests under Holm–Bonferroni correction. **Results.** Relaxing the criterion from control attribution to any finding shifts apparent recall by 6.7 to 10.0 percentage points for source- and plan-level tools, under half the 23.3-point spread between them, and by 90.0 points for repository-edge scanning, promoting it from last to second. No difference among the four survives correction; under it, three of their six pairs could not have survived under any outcome, and the rest only with a unanimous split. A third-party scanner, not our pipeline, attains the highest point-estimate recall. **Conclusion.** A comparison that does not state its matching criterion cannot be interpreted, since that choice alone moved one layer from last to second. We release our open replication package; every results table is a generated artefact.

Keywords Infrastructure as Code · Security validation · Benchmarking methodology · Static analysis · Policy as code · Empirical evaluation

M. Chittineni
Independent Researcher, London, United Kingdom
E-mail: manideep.chittineni@hotmail.com

1 Introduction

Cloud platforms are increasingly provisioned through Infrastructure as Code (IaC), in which infrastructure topology is expressed declaratively in version-controlled artefacts (Morris, 2020; Guerriero et al., 2019), atop service models whose elasticity makes such automation a practical necessity (Armbrust et al., 2010). The property that makes IaC attractive, deterministic and repeatable application of a declared state, also amplifies its failure modes: a single insecure declaration, such as an unrestricted ingress rule or a disabled encryption setting, propagates automatically to every environment that consumes the module. Empirical studies of configuration repositories confirm that such defects are common, recurring as hard-coded secrets, excessive privileges, and insecure defaults (Rahman et al., 2019b, 2021; Verdet et al., 2023, 2025). Rahman et al. (2019a) map the resulting research area and identify tooling and empirical evaluation as two of its four constituent topics.

Because remediation after provisioning implies a window of exposure, contemporary practice shifts validation earlier, into continuous integration (Bass et al., 2015; Humble and Farley, 2010). Doing so raises an evaluation problem the literature has not resolved. Candidate gates operate at incomparable abstraction levels: lexical scanning of HashiCorp Configuration Language (HCL) source, abstract-syntax-tree analysis, and policy evaluation over resolved plan state. Each reports findings in a bespoke schema, with rule identifiers that carry no shared semantics. Comparative claims are consequently difficult to reproduce, and when a benchmark is authored by the developer of one of the tools under comparison, agreement between that tool and the ground truth is close to guaranteed by construction.

This paper addresses the evaluation problem; the pipeline it evaluates alongside third-party tools composes existing mechanisms and is not a new analyser. Two commitments shape the design.

First, *normalization before comparison*. A canonical control taxonomy mediates between native rule identifiers and benchmark ground truth, so that a finding is credited when it identifies the control a case was authored to violate rather than merely because the tool emitted something. Findings whose rule identifiers the taxonomy does not cover are counted and recorded in the released results rather than discarded, since silently treating an unmapped detection as a miss biases the comparison toward whichever tool the taxonomy was authored around.

Second, *measurement separated from assumption*. A case is admitted only by a mechanically enforced rule: valid HCL2 referencing resource types the declared provider defines, a non-contradictory ground-truth label, and an unambiguous mapping onto canonical controls. The reported corpus size is the admissible count, not the number of entries in a catalogue. A tool that is not installed in the measurement environment is reported as such and excluded; it is never represented by an assumed detection rate. Latency is sampled repeatedly and reported with dispersion.

The design follows the benchmarking requirements Hasselbring (2021) sets out for software engineering research, in particular that a benchmark publish its admission rule and its measurement configuration rather than only its headline scores.

One outcome should be stated at the outset rather than discovered in Section 5. Applied to the pipeline this paper evaluates, the procedure reports a third-party scanner ahead of it on point estimates. We report that unadjusted; it is one indication, though not proof, that the procedure does not flatter the party that wrote it.

1.1 Contributions

- A finding-normalization methodology and canonical control taxonomy that render heterogeneous IaC scanner output commensurable, with three explicit matching criteria that separate correct attribution from incidental detection, and a measurement of how far the choice among them moves the conclusion. On this corpus that choice moves recall by 29 to 43 percent of the spread between the tools compared, and moves one layer from last to second.
- A statistical protocol for paired, small-sample scanner comparison: exact Clopper–Pearson intervals (Clopper and Pearson, 1934), the exact binomial form of McNemar’s test (McNemar, 1947; Dietterich, 1998), Holm–Bonferroni correction (Holm, 1979), and $d_{\min}$, the smallest discordant imbalance a comparison could have resolved. Under the correction applied here, three of the six comparisons among the four source- and plan-level tools could not have reached significance under any outcome, and the other three only under a unanimous split; $d_{\min}$ is what distinguishes that from evidence of similarity, and it is computable from the discordant count alone.
- An open replication package — corpus, generator specifications, admissibility gate, normalization engine, execution harness, and analysis code — in which every reported table is a generated artefact traceable to recorded tool output, and which retains the measurements taken before the disclosed plan-level policy correction so that the disclosure can be checked rather than taken on trust.

2 Background and Related Work

2.1 Security Weaknesses in Infrastructure as Code

Rahman et al. (2019b) established that infrastructure scripts contain recurring security smells, first across Puppet and subsequently in Ansible and Chef (Rahman et al., 2021), and derived a defect taxonomy for IaC scripts (Rahman et al., 2020). Sharma et al. (2016) had earlier catalogued configuration smells over 4,621 Puppet repositories, establishing the maintainability counterpart

to the security work. Complementary studies characterise misconfiguration in Kubernetes manifests (Rahman et al., 2023), Kubernetes security practices (Shamim et al., 2020), security practice in Terraform repositories specifically (Verdet et al., 2023, 2025), and practitioner conventions drawn from grey literature (Kumara et al., 2021). Quality and defect-prediction metrics for infrastructure code have also been catalogued (Dalla Palma et al., 2020, 2022). This body of work motivates automated pre-deployment validation but does not address how competing validation tools should be compared.

2.2 Static Analysis of IaC and Existing Comparisons

Chiari et al. (2022) survey static analysis for IaC and note the absence of shared evaluation infrastructure. The closest detector-side prior art is GLITCH, which detects security smells across several IaC technologies via an intermediate representation and is evaluated against manually annotated oracle datasets with a tool comparison (Saavedra and Ferreira, 2022; Saavedra et al., 2023). The contrast with this work is a clean one: GLITCH normalizes *inputs*, so that one detector can analyse five languages, whereas IaCSecBench normalizes *outputs*, so that detectors reporting incommensurable findings can be scored against one ground truth. Section 5.5 states why GLITCH’s oracles are not the corpus used here and what would be required to use them. The taxonomy such detectors are scored against is itself under revision, most recently beyond the original seven smell categories (War et al., 2025), which is one reason this paper versions its control set as data rather than fixing it in prose.

Comparative evaluation of the practitioner scanners has begun to appear. Roe et al. (2026) benchmark IaC scanners across multiple cloud providers and report per-tool detection counts; Gokhale et al. (2026) contrast policy-driven evaluation with taint-aware static reasoning for IaC misconfiguration; and Kapetanidou et al. (2025) evaluate the equivalent tool class for Kubernetes. These establish that the comparison is wanted. None addresses the two obstacles this paper targets: findings are compared in whichever schema each tool emits, so a rule that fires for the wrong reason is indistinguishable from one that fires for the right one, and the corpora are not released with an admissibility procedure a reader can re-run. Our contribution is orthogonal to these studies and could be applied to their corpora.

Opdebeeck et al. (2023) show that control- and data-flow analysis materially changes security-smell detection in IaC (see also Opdebeeck et al., 2022). That result motivates the plan-level layer evaluated here: a compiled Terraform plan is a resolved artefact in which interpolation, conditionals and expansion have already been applied, so evaluating it sidesteps a class of analysis difficulty instead of solving it. Qiu et al. (2024) pursue the complementary route of recovering semantic checks over IaC programs directly, and policy-violation analysis has been extended to packaged Kubernetes artefacts (Minna et al., 2026). Empirical work on policy-as-code adoption indicates which controls practitioners actually enforce (Foalem et al., 2026).

2.3 Evaluation Methodology

Methodologically this paper follows the static-analyser evaluation tradition established by Habib and Pradel (2018), who show that headline detection rates are highly sensitive to what counts as a match, and the benchmark-construction cautions of the OWASP Benchmark project (OWASP Foundation, 2023). The three-level matching criterion of Section 3.4 is a direct response to the first. Our reporting conventions are those of the static application security testing (SAST) tool-comparison literature, which reports per-category rather than pooled effectiveness and separates the absence of a rule from the failure of one (Li et al., 2023; Esposito et al., 2024). Böhme et al. (2022) show, for fuzzer benchmarking, that a benchmark’s conclusions depend on evaluation choices a report may leave unstated, which is the argument for the version manifest of Section 4.5 and for preferring interval width to corpus size as an indicator of evidential strength. Threats to validity are organised under the categories of Wohlin et al. (2012), and the study is reported against the empirical standards for software engineering research (Ralph, 2021; Hasselbring, 2021).

2.4 Policy as Code and Native Infrastructure Testing

Open Policy Agent (OPA) provides general-purpose policy evaluation over structured documents (Open Policy Agent, 2024), and Terraform provides a native test harness for module behaviour (HashiCorp, 2024). Compliance frameworks such as the Center for Internet Security (CIS) Amazon Web Services Foundations Benchmark, version 3.0.0 (Center for Internet Security, 2024), and NIST SP 800-53 (National Institute of Standards and Technology, 2020) catalogue the requirements that scanner rules are commonly mapped to; where a benchmark control has a CIS counterpart, the benchmark cites it. IaCSecBench composes these mechanisms; it does not extend them.

3 Framework Architecture

IaCSecBench composes three validation layers, each addressing a distinct failure class. Fig. 1 shows their arrangement and the gating relation between them: a commit reaches a later layer only by passing every earlier one, so the layers compose as a conjunction in execution.

Detection, however, is scored as a union, and the two are consistent. Under gating, a case that any layer would catch is a case the pipeline catches, because the commit is stopped at the first layer that objects; the conjunction governs which layers get to run, not which defects get caught. Where this paper reports the pipeline’s detection effectiveness, the figure is therefore the union of layers 1 and 3, marked as such.

Two of the three layers are scored per benchmark case. Layer 2 validates module behaviour rather than detecting resource misconfiguration, so it has




Fig. 1 Layered validation pipeline: repository-edge data validation, native module testing, and policy evaluation over compiled plan documents. A change reaches a later layer only by passing every earlier one. Edge labels give the mean per-case latency of the scored layers; the full figures, including dispersion, are in Table 7. Layer 2 validates module behaviour rather than resource configuration, so it carries no detection count on this corpus and is reported by suite size in Section 5.3.

no detection count on a corpus of misconfigured resources and is reported by suite size alone in Section 5.3.

3.1 Layer 1: Repository-Edge Data Validation

The first layer prevents sensitive material from entering version control through schema verification, detection of sensitive fields, and pattern-based identification of personal data. It operates on repository content rather than infrastructure semantics, and its characteristic failure mode is over-detection: high-entropy strings in output declarations are frequently flagged as credentials. Over-detection is not a cosmetic cost. Developers discount tools whose warnings they learn to distrust, an effect documented both for static analysis generally (Vassallo et al., 2019) and for security warnings specifically (Tahaei et al., 2021), to the point that identifying which warnings merit action is now a research area in its own right (Ge et al., 2023). This layer is therefore reported as its own row in every table of Section 5 rather than pooled with the policy layers, so its false positives stay visible.

3.2 Layer 2: Native Module Testing

The second layer validates module behaviour using Terraform’s native test framework (HashiCorp, 2024), covering variable constraints, output correctness, conditional resource creation, and tagging requirements without provisioning cloud resources.

3.3 Layer 3: Policy Evaluation over Compiled Plans

The third layer converts a Terraform plan to its JSON representation and evaluates policies written in Rego, OPA’s policy language, against it (Open Policy Agent, 2024). Because the plan is produced after expression resolution, policies observe concrete attribute values rather than unevaluated references. Plan generation is configured so that provider credential validation is disabled, which permits evaluation in continuous integration without cloud account access; this is a precondition for the layer being usable at all, and is implemented by an injected provider override, never by relaxing a policy.

3.4 Finding Normalization

Comparing heterogeneous scanners requires a canonical representation. A finding is normalised to

$$\mathcal{F} = \langle c, t, \kappa, r, s \rangle, \quad (1)$$

where c identifies the benchmark case, t the tool, κ the set of canonical controls the tool’s native rule maps to, r the resource address, and s the reported severity. The mapping from native rule to κ is many-to-many: a coarse rule may legitimately cover several controls, and such rules are reported explicitly because they weaken attribution precision.

Detection is then evaluated at three strictness levels:

- *control*: some finding’s κ intersects the controls the case was authored to violate. This is the primary criterion.
- *resource*: some finding names the resource address the ground truth identifies, irrespective of which rule fired. This isolates the correct-resource-wrong-reason case.
- *any*: the tool emitted any finding within the case. This criterion overstates detection and is reported to quantify how much of a tool’s apparent performance is incidental.

Reporting all three is deliberate: adopting a single permissive criterion is the most common mechanism by which scanner benchmarks overstate detection (Habib and Pradel, 2018). Where the ground truth does not name a resource address, the resource criterion is inapplicable, and such cases are left out of its denominator; on this corpus that excludes the four external cases.

The three are not a single ordering. Only *control* and *any* nest, because a finding that maps to the case’s control is a finding. *Resource* is a different question from either, and it departs from *control* in both directions. On the 26 vulnerable cases whose ground truth names a resource address, it credits each of Checkov, tfsec, Trivy and the plan-level layer with one case more than *control* does, and the repository-edge layer with one fewer: on those cases *control* gives 88.46, 73.08, 69.23, 61.54 and 3.85 percent and *resource* 92.31, 76.92, 73.08, 65.38 and 0.00. A tool may fire on the right resource for a different control, which *resource* credits and *control* does not, or attribute the right control to a neighbouring resource address — the bucket rather than the encryption-configuration resource the ground truth names, say — which *control* credits and *resource* does not. On this corpus the first accounts for five tool–case pairs and the second for one, and the repository-edge layer’s single control-level detection carries no resource address at all. Where this paper speaks of relaxing the criterion, and where the Δ column of Table 2 is computed, the movement is from *control* to *any* along the one axis that does nest. *Resource* is reported alongside them as a separate diagnostic of attribution precision, not as an intermediate point on that axis.

Fig. 2 traces a finding through the engine. The ordering matters for the comparison’s validity: normalization onto the canonical taxonomy precedes confusion-matrix construction, so no tool is scored in its own reporting schema and the matching criterion is applied identically to every tool.

4 Evaluation Methodology

The threat classes the pipeline’s controls were selected against, and the four domains held out of scope, are given in Appendix A. That decomposition delimits intent; the operative statement of what this evaluation measured is the corpus coverage of Section 4.2.




Fig. 2 Finding-normalization engine and evaluation pipeline: heterogeneous scanner output is mapped onto the canonical control taxonomy before any confusion matrix is formed.

4.1 Corpus Construction and Admissibility

The corpus separates a controlled collection, authored for this evaluation, from an external collection of four configurations illustrating published CIS controls. The distinction was meant to separate cases that share an author with the policy set, and so cannot confirm detection capability independently, from cases that do not. The replication package does not record an upstream source for the external configurations, however, and they entered the repository in a commit by the investigator, so we do not rely on them for independence (Section 5.5).

Admissibility is enforced mechanically. A case is admitted only if it satisfies all of the following:

1. it declares at least one resource or module block;
2. its ground-truth label is present and internally consistent, where a case declaring contradictory labels is rejected rather than resolved by precedence;
3. it resolves to exactly one interpretation in the canonical taxonomy, since a CIS section number may span several distinct controls and inference cannot choose between them;

Table 1 Corpus admissibility. The admissible count is the corpus size reported throughout; catalogue entries without configuration are not cases. The declared and present counts are given per collection because the two collections diverge by very different factors. Rejection reasons are determined mechanically by `evaluation/corpus.py`.

Corpus status	Internal	External	Total
Catalogue entries declared	345	175	520
Configurations present on disk	52	4	56
Rejected	0	0	0
Admissible	52	4	56
vulnerable	26	4	30
compliant baseline	26	0	26

4. `terraform init` and `terraform validate` succeed, which is the only check that detects a case referencing a resource type the provider does not define.

Table 1 reports the outcome, per collection. The admissible count is the corpus size used throughout; catalogue entries without corresponding configuration are not cases and are excluded. The two collections diverge by very different factors, which is why the table is split: the controlled catalogue declares 345 entries against 52 configurations, and the external catalogue 175 against 4.

On the released corpus the gate rejected nothing: every configuration present on disk was admitted. We report that rather than leaving the reader to infer a filtering effect from the presence of a filter. What the procedure buys here is not a rejection count but a stated and re-runnable admission rule, which is the property Hasselbring (2021) asks a benchmark to publish; its value as a filter is untested on a corpus this clean and would be established only by a corpus that fails it.

Both gaps require explanation, and they have different causes. The external manifests declare 50 CIS, 75 SecureFlag and 50 Terraform Registry entries, of which only four CIS configurations exist and the other two collections have none; a manifest entry without a configuration is a citation rather than a case, and counting such entries would inflate the external collection by more than a factor of forty while contributing no measurement. The controlled catalogue is the residue of a larger planned corpus: none of its 345 entries has a configuration on disk, and the 52 controlled cases were instead generated from per-control specifications under identifiers of their own. We report the declared count rather than pruning the catalogue to match, because the gap is itself the datum that motivates reporting admissible rather than catalogued size, and because the catalogue indicates where the corpus can be extended. No reported figure derives from a declared-but-absent entry.

4.2 Control Coverage

The canonical taxonomy defines 26 controls, and every one is now exercised by at least one admissible case. An earlier revision of this corpus exercised 22: four controls, one for API-gateway authorization and three for Kubernetes workloads, were declared in the taxonomy with nothing run against them. A control in that state is invisible in the results rather than reported as a gap, because an absent case produces no row, so the four were specified and generated.

Closing them established a distinction the confusion matrix cannot express. Neither tfsec nor Trivy emits any finding on any `kubernetes_*` Terraform resource — not a different finding, none at all, verified on both `kubernetes_pod` and `kubernetes_pod_security_policy` and on the compliant variants as well as the vulnerable ones. Only Checkov inspects them. Their non-detection on those three controls therefore measures the scope of what each tool reads, not its accuracy, while §5.2 scores it identically to a missed misconfiguration. The plan-level layer’s misses on the same cases have a different cause: its policy set implements no rule for the three controls, so they are a coverage gap in the reference implementation and are scored as misses. The control map records tfsec and Trivy, and only those two, as inapplicable for those controls and the analysis emits a caveat naming them, but the tables cannot separate the two failures, so we report the effect directly: including the Kubernetes controls moves control-level recall by +1.1 points for Checkov, −8.5 for tfsec, −8.1 for Trivy and −7.4 for the plan-level layer. Treating tfsec’s and Trivy’s Kubernetes cases as inapplicable rather than missed would raise their recall to 85.19 and 81.48 but leave the best-to-worst spread at 23.33 points, because its endpoints are Checkov and the plan-level layer (§5.1).

Auditing the taxonomy against the installed tool catalogues found three defect classes in the mapping itself, all corrected. Seven controls named plan-level policy rules the policy set does not implement. Two named Checkov identifiers absent from the installed catalogue entirely, and one mapped a resource-*requests* check to a resource-*limits* control, which are distinct Kubernetes concepts. None of these can produce a spurious detection, because nothing emits the identifier; each instead inflates the apparent coverage of the layer it is attributed to, and in the first case that layer is the reference implementation. That is the direction of error that flatters the contribution under evaluation, which is why the audit was mechanical, and why the resulting invariant, that every plan-level identifier must exist in the policy source, is now enforced by a test. A later audit found one defect that could have produced a spurious detection: tfsec’s and Trivy’s Lambda tracing rules were mapped to the function-encryption control. Neither fires on that control’s cases, so no classification changed, and both were removed; tfsec and Trivy now have no rule for that control, so their miss on it is an absent rule rather than a failed one.

The correction left two controls with no implementing rule in any tool, and so detectable by none of them. A control in that state scores every case citing it

as a miss for every tool simultaneously, which reads as unanimous tool failure rather than as absent coverage; both were removed, and neither was exercised, so no reported figure changes. Two controls retain an unverified flag, down from eight. The eight were retained on the grounds that this environment could not corroborate a tfsec long identifier: the recorded tfsec release provides no check-listing command, and its `--filter-results` option returns nothing for an unknown identifier and for a real but untriggered one alike. Purpose-built cases settle the question that the tool's interface cannot. An identifier that fires on the vulnerable member of a pair and stays silent on the compliant one is corroborated by observation regardless of what the catalogue can be made to list, and on that basis three of the eight were confirmed against emitted output and their evidence recorded alongside the mapping. Three more, the Kubernetes controls, were resolved rather than confirmed: tfsec emits nothing on Kubernetes resources, so its identifiers there were moved out of scope and tfsec and Trivy recorded as inapplicable, with Checkov's firing recorded as the evidence for each control.

The two survivors are comparator gaps rather than mapping errors, and we state what was observed instead. For a compute instance assigned a public address, Checkov fires but tfsec emits nothing for `associate_public_ip_address` on an instance; it flags only the subnet-level form when `map_public_ip_on_launch` is set, which this case does not set. For credential material embedded in configuration, no comparator emits a finding mapped to that control, although all three report other findings on the same resource: Checkov's secret checks belong to its secrets framework and the harness invokes it with `--framework terraform`, so they cannot fire, and only the reference pipeline's repository-edge layer detects the case. That second asymmetry favours the contribution under evaluation, which is why it is stated here rather than left in a mapping file. Identifiers are retained rather than deleted, because a wrong identifier is inert and the only possible effect of keeping it is to credit a comparator if it turns out to be real. Where the two options differ we take the one that cannot flatter the reference.

4.3 Ground-Truth Labelling

Labels are derived mechanically rather than by human rating. Each controlled case is emitted by a generator from a per-control specification that declares the intended violation, and the label follows from which member of the vulnerable/compliant pair is being written. The generator, the specification, and the emitted cases are all in the replication package, so any reader can check that a case carries the label its specification declares.

Because labels are a deterministic function of the specifications, rater disagreement is impossible by construction, and the threat that replaces it is *specification error*: a specification that misreads its originating control produces a wrong label reproduced identically across the whole pair. An agree-

ment coefficient computed over the generator’s own output would measure its determinism, not concurrence, so we do not report one.

We instead ran a blind second labelling pass against the threat that does apply. It is an *automated consistency audit*, not a second rating in the inter-rater sense: **the second labeller is a large language model, not a second human investigator**. We name it that way throughout because the word *rater* imports a reliability claim this design cannot support, and the distinction governs how the agreement below may be read.

The identity and version of the auditing model, its access route and decoding parameters, and the exact instruction it was given were not recorded when the pass was run, and we do not reconstruct them after the fact. What is recorded is the blinding procedure described below and the label and reason returned for each case, in `benchmark/labelling/independent_relabelling.json`. No case’s labelling carried another case’s label or reason by construction, but the cases were labelled one at a time within a single session, so earlier cases’ labels and reasons may have been in context when later ones were labelled, and the record cannot exclude it. Five cases are a more direct exception, discussed at the end of this section: the model had already seen their configurations and recorded labels. The omission matters: the agreement reported below is a property of one model under one instruction, and with neither identified the labelling cannot be repeated. The per-case labels are recorded, so the agreement figure can be recomputed from them, but it cannot be re-obtained from the instrument. We therefore report it as a record of what was observed, not as a reproducible measurement, and the single disagreement it surfaced stands on its own analysis below rather than on the audit’s authority. The model was used as a measuring instrument in this role; it is not an author, and responsibility for how its labels are used here rests with the author. The other uses of AI assistance in this study are described in Section 4.5.

For each admissible case a review view was constructed containing only the Terraform configuration and the canonical control under test, with every comment line removed, since the generator writes the expected label and its rationale into each file header, and with the annotation files and the case identifier withheld, the latter because its suffix names the class. Cases were presented in an order keyed on a hash of the case identifier, and the views were screened for residual label vocabulary. The recorded view contained only the configuration and the control text; the instruction that accompanied it was not recorded. A label and a reason were recorded for each case before any comparison.

Agreement is reported before reconciliation, since resolving disagreements guarantees consensus: 47 of 48 cases agreed, giving raw agreement of 97.92% and Cohen’s $\kappa = 0.958$ against a chance-agreement baseline of 50.17%. We report κ as a consistency figure for one model under one prompt and not as an inter-rater reliability coefficient, for which this design supplies no evidence. With the class balance of the audited set, 26 vulnerable against 22 compliant, expected agreement is close to one half, so κ here carries little information beyond the raw count; the 2×2 agreement table is in the replication package.

The audited set is 48 of the 56 admissible cases, and the eight it excludes are the four pairs added to close the coverage gap of §4.2. The relabelling pass was run once, blind, against the corpus as it then stood; the later cases have not been through it, and running the model again would not extend the same audit, because the original model and instruction are unrecorded and a new run would be a different instrument. We therefore report the audit on the set it covered rather than extrapolating its agreement rate to the whole corpus, and note that the eight unaudited cases are precisely the ones whose labels were written last and reviewed least.

The single disagreement is the kind of defect this pass exists to surface. A bucket configured with SSE-S3, server-side encryption with Amazon S3 managed keys (AES256), was labelled violating under a control titled *lacks server-side encryption* and citing CIS Amazon Web Services 2.1.1. That is a requirement of version 1.4.0 of the benchmark, *Ensure all S3 buckets employ encryption-at-rest* (Center for Internet Security, 2021), which version 3.0.0 removed. The configuration does not lack server-side encryption, and that requirement is satisfied by SSE-S3; read against its stated text the case is compliant. The generator’s rationale reveals the intent, namely encryption *other than* a customer-managed key, so the label encoded a requirement the control text never stated. That stricter requirement is also what the reference policy set enforces, which makes this a concrete instance of the construct-validity threat of Section 7.1 rather than a hypothetical one: the ground truth agreed with the reference implementation rather than with the control it cited.

We corrected the control text rather than the label, because requiring a customer-managed key is a defensible control and is what the corpus, the policy set, and the customer-key rules of all three source-level scanners test; relabelling instead would have left the pair testing nothing. The control is retitled to state the key requirement and records that it is stricter than the version 1.4.0 requirement it derives from; it cites no version 3.0.0 requirement, because there is none. No confusion-matrix entry changed. A regression test now pins the citation in the taxonomy to the one in the generator specification so the two cannot drift apart silently.

A later audit of every citation against version 3.0.0 (Center for Internet Security, 2024) found the taxonomy mixing its numbering with that of version 1.4.0, and several citations naming unrelated requirements. All are now given against version 3.0.0; the five controls whose citations had no counterpart there cite none, with the reason recorded in the control map. The four external configurations carry labels of their own that follow no version’s numbering, and their metadata records the corresponding version 3.0.0 requirement. No label or classification depends on a citation number, so no reported result changed.

Three properties bound what this pass establishes, and we state each at the strength the record supports. The audit reads the control texts from the same public documentation the specifications were written from, so the two readings are not fully independent and agreement is biased upward accordingly. The model was run in the same assistant session as earlier debugging, during

which it had seen the configurations and recorded labels of five cases, so their agreement is not blind; excluding all five leaves 42 of 43 agreeing, with the disagreement outside that set. And the blinding was not perfect: one external case declares `resource "aws_cloudtrail" "insecure_trail"`, a resource name that hints at the label. The screen for residual label vocabulary flagged it, and it was left unmodified rather than renamed, so that the released case is the one that was audited; the case is identified in the replication package so its contribution to the agreement figure can be discounted. The full per-case labels, reasons, and these caveats are in the replication package.

Two further checks bear on specification error. The admissibility procedure of Section 4.1 rejects any case whose label is absent, internally contradictory, or inferred to resolve to more than one control in the taxonomy; a case may still declare several controls explicitly, as one external CIS case does. And the audit of the control map reported in Section 4.2 found three classes of mapping defect by mechanical comparison against the installed tool catalogues. Neither check can establish that a specification reads its control correctly; that remains an unmitigated threat, and independent relabelling by a second investigator is the obvious next step.

4.4 Research Questions

The matching criterion is the primary object of study, and the research questions are ordered accordingly.

- **RQ1:** How much does the matching criterion change apparent detection, relative to the differences between the tools being compared?
- **RQ2:** What detection effectiveness do the evaluated tools achieve on the admissible corpus, how precisely is that effectiveness estimated, and what difference could the comparison have detected?
- **RQ3:** What execution latency does each layer contribute in continuous integration, and what is the size of the native validation suite required to express the corresponding assertions?
- **RQ4:** Which layers account for which detections?
- **RQ5:** Does effectiveness persist on independently sourced configurations?

The inferential statistics that answer RQ2 are there to characterise uncertainty and to bound what a comparison at this corpus size could resolve. They are not there to establish that one tool is superior, and no claim of superiority or of equivalence is drawn from them anywhere in this paper.

4.5 Measurement Environment and Tool Versions

Tool versions are resolved at execution time and recorded in `results/run_manifest.json`, together with the platform, processor, and repository commit; for the reported run that commit does not identify the measured tree, as

Section 7.1 sets out. Versions are not restated in prose, because a hand-copied version string is the first thing to drift; the manifest is authoritative (Böhme et al., 2022).

Tools not installed in the measurement environment are recorded as such and omitted from every table. No tool is assigned an assumed detection rate or an assumed latency.

AI coding assistants were used in building the study: to write parts of the harness and analysis code, the plan-level policy set and the benchmark cases, and to help debug the harness and inspect results, and an AI assistant was used to draft revisions of manuscript passages under the author’s direction. Which assistant products were used is stated in the Statements and Declarations. None of these uses is an instrument of the detection measurements: every reported detection and latency figure is the recorded output of the tools under test, scored by the released harness, and every artefact the assistants contributed to is in the replication package for inspection. The one instrument role a language model did play, as second labeller in the consistency audit, is described in Section 4.3. What their use in writing both the policies and the cases implies for construct validity is discussed in Section 7.1.

4.6 Tool Selection

The comparison covers Checkov, tfsec, Trivy, and standalone policy evaluation over plan documents. Selection was governed by two criteria: the tool must parse HCL2 or Terraform plan JSON without requiring cloud provider API access, and the set must span distinct analysis paradigms, namely syntax-tree analysis, lexical scanning, and policy evaluation. KICS, Regula, and provider-native configuration services are excluded under the first criterion or as duplicating a represented paradigm.

Trivy requires separate comment, because it is in the set for a reason unrelated to paradigm coverage. Aqua Security folded tfsec’s engine and rule set into Trivy, and Trivy is the maintained product; tfsec is no longer developed. A comparison reporting tfsec alone would be characterising a scanner practitioners are advised off, so both are measured. But the two are not independent implementations. Their rule identifiers are the same Aqua Vulnerability Database numbers under different spellings (tfsec reports `AVD-AWS-0086` where Trivy reports `AWS-0086`), and we exploit that correspondence to build Trivy’s entries in the control taxonomy: each is derived mechanically from the identifier pair tfsec itself emits, rather than chosen by observing which rules Trivy happens to fire on a case, which would fit the taxonomy to the tool. A regression test re-derives the crosswalk and fails if the two disagree.

The consequence bounds what the enlarged tool set establishes. Three third-party scanners are measured, but only *two* independent third-party rule sets: Checkov’s, and the tfsec/Trivy lineage. Agreement between tfsec and Trivy is therefore close to a consistency check on that shared lineage across two engine generations rather than corroboration by a second party, and we do

Table 2 Recall under three finding-matching criteria. *Control* credits a detection only when the tool’s rule maps to the control the case violates; *resource* credits any finding on the expected resource; *any* credits any finding whatsoever. The criteria are not a single ordering. *Resource* is computed only over the vulnerable cases whose ground truth names a resource address (26 of 30 here), so its column is not directly comparable with the other two. Δ is the control-to-any spread, reported as the count of vulnerable cases the permissive criterion adds and as the corresponding percentage with an exact Clopper–Pearson interval on that count. It is given per tool rather than summarised, because the spread differs by an order of magnitude across the set and a single headline figure would be carried entirely by its largest member.

Tool	Control (%)	Resource (%)	Any (%)	Δ (pts, 95% CI, cases)
Checkov	90.00	92.31	96.67	6.67 [0.82, 22.07] (2/30)
IaCSecBench L1	3.33	0.00	93.33	90.00 [73.47, 97.89] (27/30)
OPA (plan-level)	66.67	65.38	73.33	6.67 [0.82, 22.07] (2/30)
tfsec	76.67	76.92	83.33	6.67 [0.82, 22.07] (2/30)
Trivy	73.33	73.08	83.33	10.00 [2.11, 26.53] (3/30)

not report it as the latter. Where the two diverge, the difference is attributable to engine and rule evolution between the products, and is informative for that reason.

Terrascan is absent for a reason that is not a selection criterion, and we state it plainly: no successful Terrascan execution was ever recorded against this corpus, so there is nothing to report. Its paradigm, Rego policy evaluation, is represented in the tool set by the plan-level layer, which bounds what the omission costs without eliminating it. Policy evaluation is therefore represented only over compiled plan documents, and no source-level policy engine appears in the tool set. A source-level Rego engine could differ from a plan-level one on configurations whose violations are resolved during planning, and this evaluation does not measure that difference. We make no claim of any kind about Terrascan’s detection behaviour.

5 Results

All tables in this section are generated by `evaluation/analyze.py` from recorded raw tool output, regenerated by `experiments/run_baselines.sh`, and not edited by hand. In them, *IaCSecBench L1* is the repository-edge layer and *OPA (plan-level)* is the plan-level layer, which evaluates the author’s policy set.

5.1 RQ1: Sensitivity to the Matching Criterion

Table 2 reports recall under each criterion, with the control-to-permissive spread Δ given per tool, in cases as well as in points, and with an exact interval on the case count. Two readings of that column matter, and they differ in magnitude by an order of magnitude.

Among the three source-level scanners and the plan-level layer, Δ ranges from 6.67 to 10.00 percentage points. The spread between those same four tools at the control criterion is 23.33 points, from Checkov’s 90.00 to the plan-level layer’s 66.67, which is more than twice the criterion effect.

That comparison needs one qualification, which §4.2 sets out. Three of the 26 controls govern Kubernetes workloads. tfsec and Trivy read none of the resources those controls are expressed in, so their zero on those cases is a statement about scope rather than detection; the plan-level layer’s zero on the same cases is a gap in its own policy set and a genuine miss. Treating tfsec’s and Trivy’s Kubernetes cases as inapplicable leaves the spread at 23.33 points, because its endpoints are Checkov and the plan-level layer. Excluding the three controls altogether would narrow it to 14.81, but only by removing the reference implementation’s own missing rules, which flatters the reference, so we do not treat that figure as scope-controlled. On the point estimates the criterion effect is therefore 29 to 43 percent of the between-tool spread, and the spread is 2.3 to 3.5 times the criterion effect; with intervals as wide as those below, this corpus does not establish that ordering beyond the point estimates.

Both effects remain small in cases: two or three of the 30 vulnerable cases move for each of the four source- and plan-level tools when the criterion is relaxed. The intervals in Table 2 are correspondingly wide, so what this corpus establishes is that the criterion effect is not negligible beside the between-tool effect, not the size of either. We state the claim at that strength and no further.

Rank order is where the criterion’s effect is uneven, and it is worth being precise about where it does and does not disturb it. Among the three source-level scanners and the plan-level layer the criterion reorders nothing: Checkov leads at both criteria, the plan-level layer trails at both, and relaxation only collapses tfsec and Trivy into a tie at 83.33. The reordering is entirely the repository-edge layer’s, and it is large. That layer’s Δ is 90.00 points, from 3.33 at the control criterion to 93.33 under the permissive one, which moves it from last of the five to second. It detects hardcoded credentials and comparable textual patterns, so on a corpus dominated by attribute-level controls it identifies the seeded violation in exactly one case while emitting some finding in almost all of them. The ratio between its two figures is close to twenty-eight, and we report the difference in points rather than that ratio: with one true positive in the denominator the ratio is an artefact of a small count, whereas the 90-point difference is the quantity that would mislead a reader of a single-criterion report.

The two readings support different claims, and only the second is strong here. A single-criterion report does not necessarily misstate the *ranking* of tools working at the same abstraction level; on this corpus it does not. It can misstate the standing of a tool whose findings are of a different kind by most of the available range, and it gives a reader no way to tell which of the two situations they are in. That is the argument for reporting the criterion, and it does not require that the criterion have reordered anything.

Table 3 Confusion-matrix counts on the admissible corpus ($N = 56$: 30 vulnerable, 26 compliant). Rates and intervals over the same counts are in Table 4.

Tool	TP	FP	TN	FN	MCC
Checkov	27	0	26	3	0.898
IaCSecBench L1	1	1	25	29	-0.014
OPA (plan-level)	20	0	26	10	0.694
tfsec	23	2	24	7	0.692
Trivy	22	1	25	8	0.704

Table 4 Detection rates with exact Clopper–Pearson 95% intervals, over the counts of Table 3. Recall’s denominator is the vulnerable-case count fixed by the corpus design; precision’s is the number of findings the tool emitted, so its interval is conditional on that realised count rather than a marginal statement about the tool. Entries marked n/e are not estimable from this corpus.

Tool	Recall %	Precision %	Specificity %
Checkov	90.00 [73.47, 97.89]	100.00 [87.23, 100.00]	100.00 [86.77, 100.00]
IaCSecBench L1	3.33 [0.08, 17.22]	50.00 [1.26, 98.74]	96.15 [80.36, 99.90]
OPA (plan-level)	66.67 [47.19, 82.71]	100.00 [83.16, 100.00]	100.00 [86.77, 100.00]
tfsec	76.67 [57.72, 90.07]	92.00 [73.97, 99.02]	92.31 [74.87, 99.05]
Trivy	73.33 [54.11, 87.72]	95.65 [78.05, 99.89]	96.15 [80.36, 99.90]

In answer to RQ1, relaxing the criterion from *control* to *any* moves recall by 6.67 to 10.00 points for the source- and plan-level tools, 29 to 43 percent of the 23.33-point spread between them, and by 90.00 points for the repository-edge layer, which it moves from last to second.

5.2 RQ2: Detection Effectiveness and Resolvable Differences

Table 3 reports confusion-matrix counts and Table 4 the corresponding rates with exact Clopper–Pearson intervals. Intervals accompany every proportion because a point estimate on a corpus of this size is compatible with a wide range of true values; Checkov and the plan-level layer, attaining 100.00% precision, carry lower confidence bounds of 87.23 and 83.16 respectively, and reporting the point estimate alone would misstate what has been established.

The Clopper–Pearson form is chosen for its coverage guarantee rather than for interval width. Its actual coverage is never below the nominal level, whereas the score and adjusted-Wald intervals attain nominal coverage only on average and undercover for proportions near the boundary (Agresti and Coull, 1998), which is where most of these estimates sit; Agresti and Coull draw the opposite recommendation from the same observation, preferring the approximate intervals for their width. The cost is conservatism, so every interval in Table 4 is wider than a score interval on the same counts; for a claim about detection capability that is the safe direction of error, and the harness computes the score interval as well, so a reader preferring it can obtain it from the released results.

Table 5 Exact McNemar comparisons against OPA (plan-level). b counts cases the reference classifies correctly and the comparator does not, counting both missed violations and spurious findings; c is the converse. p is the two-sided exact binomial value, p_{adj} its Holm–Bonferroni correction across the tool family. The odds ratio carries a Haldane–Anscombe correction so that a zero discordant cell yields a finite estimate. d_{min} is the smallest $|b - c|$ the exact test could have called significant at this many discordant pairs; a dash means no split could, so the comparison had no power and its non-significance is not evidence of similarity.

Comparator	b	c	d_{min}	p	p_{adj}	OR	Cohen’s g
Checkov	2	9	9	0.065	0.196	0.26	0.318
IaCSECBench L1	21	1	12	< 0.001	< 0.001	14.33	0.455
tfsec	3	4	7	1.000	1.000	0.78	0.071
Trivy	3	4	7	1.000	1.000	0.78	0.071

Table 6 Exact McNemar comparisons over every tool pair, with Holm–Bonferroni applied across the full set of pairs rather than across one tool’s family. b counts cases the left tool classifies correctly and the right does not; c is the converse. Both error types count. The correction is stronger here than in Table 5 because the family is larger. d_{min} is the smallest $|b - c|$ this pair’s discordant count could have resolved; a dash means none could.

Tool A	Tool B	b	c	d_{min}	p	p_{adj}
Checkov	IaCSECBench L1	28	1	13	< 0.001	< 0.001
Checkov	OPA (plan-level)	9	2	9	0.065	0.393
Checkov	tfsec	7	1	8	0.070	0.393
Checkov	Trivy	7	1	8	0.070	0.393
IaCSECBench L1	OPA (plan-level)	1	21	12	< 0.001	< 0.001
IaCSECBench L1	tfsec	3	24	13	< 0.001	< 0.001
IaCSECBench L1	Trivy	2	23	11	< 0.001	< 0.001
OPA (plan-level)	tfsec	3	4	7	1.000	1.000
OPA (plan-level)	Trivy	3	4	7	1.000	1.000
tfsec	Trivy	1	1	—	1.000	1.000

One property of the precision intervals should be read with care. Recall’s denominator is the 30 vulnerable cases, fixed by the corpus design, so a binomial interval applies directly. Precision’s denominator is the number of findings the tool emitted, which is a realised outcome rather than a design parameter, so its interval is conditional on the observed prediction count and is not a marginal statement about the tool.

Specificity is estimable here because the corpus carries 26 ground-truth compliant cases, and it is reported for that reason: a detector’s cost to its users is carried as much by its false positives as by its misses, and a table reporting only recall and precision would have described the half of the behaviour that flatters a detector. The Matthews correlation coefficient (MCC) (Matthews, 1975) is likewise estimable and is given in Table 3. What the compliant baseline does not yet support is a *precise* false-positive estimate: with 26 negatives the specificity intervals span more than thirteen points even where no false positive occurred.

Tables 5 and 6 report pairwise comparisons. Table 6 is the confirmatory analysis: it covers all ten unordered pairs with Holm–Bonferroni applied across

the whole set, so the comparison a practitioner most wants, between two third-party scanners, is made, and the pipeline evaluated here does not sit on one side of every test. Table 5 retains the against-one-reference form for continuity with the way this literature usually reports, and its weaker correction reflects its smaller family; where the two disagree, the all-pairs adjustment governs.

The discordant count b includes every case in which the first tool is correct and the second is not, counting spurious findings as well as missed violations; restricting b to missed violations understates disagreement. The exact binomial p -value is reported in preference to the chi-square approximation, which is unreliable when a discordant cell is small; on these counts, where the smallest discordant cell is one, the approximation is not merely conservative but invalid. Test selection follows established guidance for software engineering experiments (Arcuri and Briand, 2014). Odds ratios carry a Haldane–Anscombe correction, so a zero discordant cell yields a finite estimate; an unbounded ratio is a degenerate cell rather than a large effect, and Cohen’s g is reported as the effect size.

No pairwise difference among the three third-party scanners and the plan-level layer is significant after correction. The $d_{\min}$ column states what each comparison could have detected, and it is the more informative result. For Checkov against the plan-level layer, eleven discordant cases mean that only a split of ten to one or more extreme would have reached $\alpha = 0.05$; the observed split was nine to two. Checkov against tfsec and against Trivy each have eight discordant cases, and the plan-level layer against tfsec and against Trivy each have seven, so each of those four comparisons could have reached significance only if every discordant case had fallen the same way. For tfsec against Trivy the discordant count is two, at which *no* split reaches significance: the smallest attainable two-sided p -value at two discordant pairs is 0.5. That comparison had no power at all, and the other five had power only against near-unanimous disagreement, so their non-significance carries little information about the tools.

Those thresholds are uncorrected, and the correction that governs is stricter. The four repository-edge comparisons occupy the first four Holm ranks, so any other pair needs $p \leq 0.05/6 \approx 0.0083$ to survive. Holding the other comparisons as observed, the plan-level layer against tfsec and against Trivy could not have survived under any split of their seven discordant cases, whose smallest attainable p is 0.0156, nor could tfsec against Trivy; the remaining three pairs could have survived only under a unanimous split. Table 6 reports $d_{\min}$ at the uncorrected level.

The one significant family is the repository-edge layer against each of the other four, which reflects that it addresses a different class of defect entirely; that is a statement about scope, not quality. The interval widths in Table 4 and the $d_{\min}$ column of Table 6 together are what this corpus establishes about relative effectiveness among the scanners, which is little, in a form a reader can act on.

In answer to RQ2, control-level recall among the four source- and plan-level tools ranges from 66.67 for the plan-level layer to 90.00 for Checkov, with

95 percent intervals between 24 and 36 points wide, and no difference among them survives correction.

5.2.1 Divergence within the shared rule lineage

The tfsec/Trivy pair (Section 4.6) permits one measurement the rest of the comparison cannot make: how much a rule set changes as it is carried from a retired product into its successor. The two agree on 54 of 56 cases. Both disagreements concern the same control, identity and access management (IAM) policy wildcards: tfsec reports `aws-iam-no-policy-wildcards` on both members of the pair, and Trivy emits its counterpart `AWS-0057` on neither, not on these two cases and nowhere in the corpus.

The divergence is not a straightforward regression. tfsec’s rule fires on the violating case, earning a true positive, *and* on the compliant case, earning a false positive; Trivy’s silence forfeits both. tfsec therefore records the higher recall and Trivy the higher precision, from a single rule behaving differently. We verified that Trivy’s miss is genuine and not an artefact of our taxonomy: the one unmapped finding Trivy emitted on that case is an S3 versioning rule, unrelated to IAM policy wildcards, so no mapping gap is concealing a detection.

The remainder of the two tools’ output corroborates the shared provenance quantitatively. Every rule Trivy emitted has a tfsec counterpart reported on an identical number of cases. The converse fails for four rules, which tfsec emits and Trivy emits nowhere in the corpus: the IAM policy-wildcard rule above; a CloudFront TLS-policy rule the taxonomy does not cover; and two mapped rules, S3 bucket encryption (`AWS-0088`, on 14 cases) and public database access (`AWS-0082`, on one), neither of which changes a classification, since the two tools disagree only on the IAM pair. Two products separated by a change of engine thus differ, on this corpus, by four rules, in every case by the successor being silent where the predecessor fires. Section 6 takes up what that implies for practitioners migrating between them.

5.3 RQ3: Execution Cost and Suite Size

Table 7 reports per-case latency as a mean with standard deviation over 168 samples per tool, three repetitions across the 56 cases. Single-run figures are not reported, because a single sample carries no estimate of its own dispersion.

On engineering effort we report a measurement of the native suite alone and withdraw the comparative claim. The validation suite comprises 11 `.tftest.hcl` files totalling 190 physical lines, of which 153 are non-blank and 137 are both non-blank and non-comment, expressing 11 `run` blocks and 26 `assert` blocks. We give both line counts because they differ by the 16 comment lines and one of them is the figure a reader would compute; quoting the larger under the stricter description would inflate the suite by 12%. Counts were obtained by line-oriented pattern matching over the suite

Table 7 Measured per-case execution latency over 3 repetitions. Values are mean $\pm$ standard deviation on the environment recorded in `results/run_manifest.json`. Each tool is measured at its own interface: the plan-level figure covers policy evaluation over an already-compiled plan document and excludes the `terraform init` and `terraform plan` invocations that produce it, which are expected to dominate that layer’s wall-clock cost in continuous integration. It is therefore a lower bound, and is not comparable with the source-level figures, which have no equivalent preparation step.

Tool	Evaluation layer	Latency (ms)
Checkov	AST static analysis	1913.58 $\pm$ 60.57
IaCSecBench L1	Repository-edge scanning	0.15 $\pm$ 0.05
OPA (plan-level)	Rego over compiled plan	12.82 $\pm$ 2.16
tfsec	HCL lexical scanning	275.22 $\pm$ 7.58
Trivy	HCL scanning, tfsec successor	400.71 $\pm$ 26.67

Table 8 Layer attribution over the vulnerable cases. Overlap is reported explicitly: layers are not assumed to detect disjoint sets. Layer 2 (native module testing) is not scored: it validates module behaviour rather than detecting resource misconfiguration, so it has no detection count on this corpus.

Configuration	Detected	Recall (%)
L1 (repository edge)	1	3.33
L3 (compiled plan)	20	66.67
L1 $\cap$ L3 (overlap)	0	0.00
Union of scored layers	21	70.00

sources, which are in the replication package so the count can be reproduced or disputed.

We report no code-volume ratio against an external test framework. That would require an independently authored equivalent suite, and constructing one ourselves would place both sides of the comparison in the hands of the party with an interest in the outcome. The claim that native testing reduces test-suite volume is therefore outside this paper’s scope. The native suite does require no language runtime, dependency manifest, or build step beyond the Terraform binary already needed to plan the configuration — a difference in toolchain surface, not a measured difference in effort.

In answer to RQ3, per-case latency ranges from 0.15 ms for the repository-edge layer to about 1.9 s for Checkov (Table 7), the plan-level figure being a lower bound that excludes plan generation, and the native suite expresses its assertions in 26 `assert` blocks across 11 files.

5.4 RQ4: Layer Attribution

Table 8 attributes detections to the layer that produced them. Each layer is scored by the same normalization and matching procedure applied to the external scanners, so a layer is credited only when its finding maps to the control the case seeds; no layer receives credit by construction.

The table reports pairwise intersections and the union alongside the per-layer counts, because the per-layer counts alone cannot distinguish a complementary pipeline from a redundant one. Reported additively, a case detected by every layer would be counted repeatedly and the composed pipeline would appear to achieve more than it does. The union is therefore the only figure that characterises the pipeline as a whole, and any excess of the summed per-layer counts over the union is redundancy rather than additional coverage.

Two structural constraints bound what this attribution can show, and both suppress apparent complementarity. The repository-edge layer is insensitive to every control expressed as a resource attribute, so on this corpus its contribution is necessarily small, as Section 5.1 already showed from the other direction. The plan-level layer can only be credited on cases for which a plan was produced; cases whose planning failed would be excluded from its denominator rather than recorded as misses. On this corpus no case was excluded on that ground, so the layer’s denominator is the full 30 vulnerable cases and the constraint is stated as a property of the attribution rule rather than as an adjustment applied here. The native module-testing layer validates the project’s own infrastructure and is not exercised per benchmark case, so it is excluded from this table and reported under RQ3; its omission is a scope boundary, not a null result.

In answer to RQ4, the plan-level layer accounts for 20 of the scored pipeline’s 21 detections and the repository-edge layer for the remaining one, with no case detected by both.

5.5 RQ5: External Generalization

Two subsets lie outside the controlled corpus, and neither supports an independent effectiveness estimate. The first is four labelled configurations illustrating CIS Amazon Web Services controls (Table 1), whose labels derive from the control each illustrates; they enter the admissible corpus on the same terms as the internal cases and are validated by the same procedure. The replication package does not record an upstream source for them, and they entered the repository in a commit by the investigator, so we do not treat them as independent of the investigator.

One defect in how they were stored is worth reporting, because its signature is invisible in aggregate. A case is scanned as a directory rather than as a file, and these four configurations initially shared one directory. Every tool therefore received identical input for all four cases and returned a byte-identical finding set for each, while the scoring stage compared that one result against four different controls. Four rows entered each confusion matrix where one observation existed. This does not add noise so much as misstate precision: interval width is driven by the number of independent observations, so duplicated cases narrow every interval that includes them without contributing evidence, and nothing in the aggregate output distinguishes this from four genuinely independent cases that happen to agree. Each configuration now

Table 9 Scanner behaviour on the unlabelled external subset: 25 third-party AWS-published Terraform repositories (15,226 lines), each pinned to an exact commit. These configurations carry no ground-truth label, so the columns count *alerts*, not detections: no value here is a true-positive or a recall figure. ‘Unmapped’ is the share of a tool’s alerts whose native rule has no entry in the canonical control map, and ‘In module’ the share reported below the repository root.

Tool	Alerts	Per kLOC	Unmapped	In module
checkov	264	17.3	73%	67%
tfsec	118	7.8	19%	47%
trivy	141	9.3	35%	75%

Table 10 Pairwise agreement between the source-level scanners on the unlabelled external subset, as Jaccard indices over the repositories both tools scanned. ‘By control’ compares canonical controls and so covers only mapped alerts; ‘by resource’ compares the Terraform resource each tool flagged and needs no control map. Agreement is not correctness: two tools may corroborate one another and both be wrong. Note that Trivy is tfsec’s successor and inherits its rule set, so that pair is not an independent comparison.

Pair	By control	By resource
checkov vs. tfsec	0.155	0.211
checkov vs. trivy	0.458	0.424
tfsec vs. trivy	0.324	0.491

occupies its own directory, and the corpus loader refuses to load a corpus in which two cases share one. The same trap applies to any evaluation whose unit of analysis is a directory: case isolation has to be enforced rather than assumed.

Four cases are too few for a per-subset recall estimate, whose interval would span most of the unit interval, so we report them within the combined admissible corpus.

The second subset is 25 Terraform repositories published by Amazon Web Services under the MIT-0 licence, 15,226 lines in all, each pinned to an exact commit, with the selection criteria recorded in the replication package. These are the only configurations here whose third-party provenance is recorded, but they carry no ground-truth label, so they yield no detection figure and nothing measured on them enters a confusion matrix: Table 9 counts alerts instead. They bear on two points made elsewhere. The share of alerts whose native rule has no entry in the control map is 73 percent for Checkov, 19 for tfsec and 35 for Trivy, against 72, 25 and 27 percent on the labelled corpus. Checkov’s share is nearly the same on both, tfsec’s is lower on production-shaped code and Trivy’s higher, so the labelled corpus does not overstate how much of a tool’s output a control-level comparison can score. And the tools agree little (Table 10): by flagged resource the Jaccard index is 0.21 for Checkov against tfsec, 0.42 against Trivy, and 0.49 between tfsec and Trivy, whose shared lineage (Section 4.6) makes the last not an independent comparison. Agreement is not correctness, since two tools may agree and both be wrong.

In answer to RQ5, this study cannot say whether effectiveness persists on independently sourced configurations: the labelled external cases are too few and of unrecorded provenance, and the independently sourced repositories are unlabelled.

We make no claim that effectiveness measured here transfers to production infrastructure. Establishing that requires a labelled corpus of real-world configurations, and the position of this literature deserves to be stated precisely rather than as an absence. Manually annotated oracle datasets of real-world IaC scripts do exist: GLITCH releases three, for Ansible, Chef and Puppet (Saavedra and Ferreira, 2022), and later versions of the tool add Terraform support (Saavedra et al., 2023). They are not the corpus used here, first because none of the oracles is a Terraform corpus, and second because their labels are security-smell categories in the sense of Rahman et al. (2019b), whereas this evaluation’s unit of ground truth is a CIS control identifier, and the two taxonomies are not in correspondence: a single smell category spans several CIS controls and several controls have no smell counterpart. Scoring the tools evaluated here against those oracles would require a mapping between the two taxonomies that does not exist and whose construction is a research contribution in itself. We identify that mapping, rather than the absence of labelled real-world data, as the principal obstacle to field-representative evaluation in this area, and return to it in Section 9.

6 Discussion

Five findings from building the harness bear on how IaC security tooling should be evaluated.

First, the matching criterion is a consequential choice in its own right. Because native rule identifiers carry no shared semantics, any comparison implicitly chooses what counts as a correct detection, and on this corpus that choice moves recall by 29 to 43 percent of the spread between the tools compared, and by most of the available range for one layer (Section 5.1). Habib and Pradel (2018) established the sensitivity for analysers of one language, which share a vocabulary of defect kinds even where they disagree about instances. The IaC case is the harder one: the tools compared here work at different abstraction levels, from lexical scanning to policy evaluation over resolved plans, and there is no shared rule vocabulary in which to state what a correct detection would be. The criterion is therefore not a tuning parameter of the comparison but a construction of its dependent variable, and it belongs in a benchmark’s reported method alongside its corpus size. The per-category reporting discipline the SAST literature arrived at is the same response to the same problem (Li et al., 2023).

Second, taxonomy coverage is a bias mechanism. When a canonical taxonomy is authored around one tool’s rule set, findings that other tools emit under identifiers the taxonomy lacks are scored as misses. The effect, where it occurs, is systematic and favours the authoring tool. Auditing observed rule

identifiers against the taxonomy after execution, and reporting the unmapped remainder, is the minimum control. On this corpus 367 findings fell outside the control map, 255 of Checkov’s, 57 of tfsec’s and 55 of Trivy’s, but only 11 of them sit on vulnerable cases the emitting tool missed (8, 1 and 2 respectively). They fall on five tool–case pairs, so at most five misses could reflect a gap in the map rather than in the tool. During this work the audit identified rules from comparator tools that the initial taxonomy omitted and that would otherwise have been scored as failures.

A third observation concerns statistical reporting rather than IaC. On corpora of this size a non-significant paired test is routinely read as evidence of similarity, yet under the correction applied here three of the six comparisons among the four source- and plan-level tools could not have been significant under any outcome, and the other three only under a unanimous split. The quantity that distinguishes the two readings, the smallest discordant imbalance the test could resolve, is computable from the discordant count alone and costs nothing to report. We suggest it accompany every paired comparison on a small corpus.

Fourth, plan-level evaluation is a trade-off rather than an improvement, and one term of that trade-off is a defect class rather than a cost. Because the compiled plan is produced after expression resolution, policies observe concrete values and avoid a class of interpolation-related imprecision that source-level analysis must otherwise reason about (Opdebeeck et al., 2023). Plan generation in exchange adds a step that the latency figures here exclude and that is expected to dominate the layer’s cost, and it requires a working provider configuration. It also introduces a distinction that a policy author must handle explicitly, because Terraform serialises three configuration states of a scalar optional attribute into shapes that a presence test orders wrongly. An attribute set to a literal carries its value. An attribute set from a reference to another resource is absent from the resolved attribute set and recorded as not-yet-known, since the value is not resolvable until apply. An attribute not set at all is present with the value null. A test for the attribute’s presence therefore passes the state that sets nothing and flags the state that sets it by reference, which does not weaken such a control but inverts it: the compliant case is reported and the violating case is let through. The two errors conceal each other in the aggregate, since the resulting confusion matrix shows one false positive and one false negative — a signature indistinguishable from ordinary imprecision. We encountered this in our own policy set and disclose the correction in Section 7.2. It is a property of plan-level evaluation rather than of that implementation: any engine reading plan documents inherits the three-way distinction, and any that collapses it inverts the same class of control. The correct treatment is to report only the not-set state and to yield no finding on the not-yet-known one, which is genuinely indeterminate at plan time. Block-typed attributes are exempt, since an unconfigured block is likewise absent but a configured one appears with its contents, so presence discriminates correctly over blocks.

A fifth observation concerns tool succession rather than tool comparison. The tfsec/Trivy pair measured here (Section 5.2) is a rule set carried from a retired product into its maintained successor across a change of engine, and on this corpus the two differ by four rules, all present in the predecessor and silent in the successor; one of them costs the successor a detection. The magnitude is small; the direction is what matters, since coverage was lost rather than gained and nothing in either product’s output announces it. Practitioners migrating between a scanner and its replacement should not assume that coverage transfers rule for rule, and a benchmark measuring only one of the two has no way to see the difference at all.

7 Threats to Validity

Threats are organised under the categories of Wohlin et al. (2012).

7.1 Construct Validity

The controlled corpus was authored by the same investigator who implemented the evaluated policy set. Ground-truth labels therefore encode the same reading of each control that the policies encode, and agreement between the two is not independent evidence of detection capability. AI coding assistance was used in writing both the policy set and the benchmark cases, so the two share that influence as well as an author, and the dependence is if anything stronger than single authorship alone implies. This is the most serious limitation of the evaluation, and it is only partly mitigated.

Two mitigations apply. Labels are intended to derive from the text of the originating control rather than from any policy implementation, though Section 4.3 documents one case where a label encoded the reference policy’s stricter reading instead, and both the specifications and the generator that consumes them are released, so a reader can audit the derivation. And the harness is also run on 25 third-party repositories the investigator did not author, though these are unlabelled and support no effectiveness estimate (Section 5.5).

A third check bears on it partially. The blind relabelling pass of Section 4.3 agreed with 47 of the 48 labels it covered, and the one disagreement it surfaced was a genuine specification defect of exactly the predicted kind: a control whose stated text was weaker than the requirement its label encoded, where the stricter reading matched the reference policy set. That the pass found such a case is evidence the mechanism is real rather than merely conceivable.

The audit is not a sufficient mitigation. Its second labeller is a language model whose reading of control texts derives from the same public documentation the specifications were written from, so the two readings share provenance, agreement is biased upward, and no claim of human inter-rater reliability is available from it. It also ran in the same assistant session used for earlier debugging (Section 4.3), so it is not independent of the AI assistance used in

building the study (Section 4.5), which weakens its independence further. Its identity, parameters and instruction were not recorded, so the audit cannot be repeated as run. A single investigator, with AI assistance, still wrote every specification, and no labelled part of the corpus has a recorded source outside this failure mode.

What would remove this threat rather than bound it is a separation of benchmark authorship from policy implementation, and we state it as a protocol because this corpus does not satisfy it: controls selected by someone other than the author of the detection rules; cases generated without reference to any detection implementation; labels fixed and recorded before the rules that detect them are written; and the corpus frozen, by commit, before any tool is executed against it. None holds fully here. The run manifest records the commit checked out at measurement time, but the reported run was taken on a working tree with uncommitted changes: the recorded commit is release 1.3.1, whose corpus had 44 internal cases, not the 52 measured. The measured corpus is therefore identified by the released package rather than by that commit, and the harness now records the tree's uncommitted paths alongside it. The first three are the substantive ones, and independent relabelling by a second human who has not read the policy set is the cheapest step toward them. A benchmark satisfying all four could offer its ground truth as independent evidence of detection capability; this one offers it as a stated and auditable assumption.

Results on the controlled corpus should therefore not be read as an estimate of field performance. A high score on a developer-authored corpus is at least as consistent with insufficient corpus difficulty as with strong detection, and we do not interpret it as the latter.

7.2 Internal Validity

Comparator tools ran with default rule sets, whereas the plan-level policies were authored with knowledge of the corpus taxonomy. This asymmetry favours the plan-level layer. It is disclosed rather than adjusted for, because equivalent tuning effort across heterogeneous engines cannot be operationalised objectively; readers should treat the comparison as bounding what default-configuration scanning achieves, not as a ranking of the tools' attainable capability.

One plan-level policy was corrected after its results had been inspected, and the correction is disclosed because it is exactly the iteration the comparators did not receive. The audit-log-encryption rule tested for the presence of a key-identifier attribute in the plan document, which inverts the control for the reason set out in Section 6: the compliant case was reported as violating and the violating case as compliant, and the two errors concealed each other in the aggregate. The corrected rule distinguishes not-set from not-yet-known and reports only the former.

Because this correction was informed by inspecting corpus results, the plan-level figures reflect one round of iteration that Checkov, tfsec and Trivy did not receive. The pre-correction and post-correction results are both in the replication package.

Latency was measured on a single host with recorded specifications, so absolute figures do not transfer across environments. Latency additionally measures each tool at its own interface: the plan-level figure covers policy evaluation over an already-compiled plan document and excludes the `terraform init` and `terraform plan` invocations that produce it, which are expected to dominate that layer’s true wall-clock cost in continuous integration. The source-level scanners have no comparable preparation step. Plan-level and source-level latencies are therefore not interchangeable, and the plan-level figure is a lower bound on deployed cost.

7.3 External Validity

The corpus targets Terraform against a single cloud provider, and the native module-testing layer is exercised against one project’s infrastructure. Neither the reported effectiveness nor the suite-size measurement of Section 5.3 generalises to other IaC technologies, other providers, or the multi-module configurations typical of large organisations. No claim is made about organisational factors such as policy ownership, escalation workflow, or developer friction, because this study measures tool output and does not observe practitioners.

The tool set bounds generalisation further: three third-party scanners are measured but only two independent rule sets are represented (Section 4.6), so the comparison speaks to those rule sets rather than to source-level IaC scanning in general. Provider coverage bounds it again, since security-policy adoption in Terraform differs materially across cloud providers (Verdet et al., 2025) and this corpus targets one.

7.4 Conclusion Validity

The corpus is deterministic, not adversarial, so reported effectiveness characterises performance on a published taxonomy of known misconfiguration classes; no claim is made regarding novel or deliberately obfuscated insecure configurations, or an adversary who has read the policy set. On statistical conclusion validity, the $d_{\min}$ column of Table 6 bounds what the paired tests could establish: under the Holm correction three of the six comparisons among the four source- and plan-level tools could not have reached significance under any outcome, and the other three only under a unanimous split. We therefore draw no conclusion of equivalence anywhere in this paper, and the effectiveness comparison among scanners should be read as unresolved rather than as settled in either direction.

8 Replication Package

The replication package contains the benchmark corpus with ground-truth annotations and the generator specifications from which those annotations derive; the canonical control taxonomy as auditable data; the per-case output of the blind relabelling pass of Section 4.3, with its reasons and its stated limitations; the execution harness, which records raw tool output, resolved versions, and repeated latency samples; the normalization engine; the statistical implementation, with a test suite pinning each method to independently derived values; and the analysis stage that generates every results table in this paper. It also retains the measured results from before the plan-level policy correction disclosed in Section 7.2, so that the disclosure can be checked rather than taken on trust.

Three properties are worth stating explicitly. The harness declines to emit results when no case is admissible, instead of producing an empty but plausible table. The statistical implementation refuses to report a proportion whose denominator is zero, returning a not-estimable marker instead of a value a reader could mistake for a measurement. And both diagrams in this paper are generated from the recorded results rather than drawn, with a continuous-integration check that fails if a re-measurement moves a corpus size or tool version that a figure asserts.

Fig. 3 gives the package layout. Ground-truth labels, the scenario corpus, and the policy sources are separate top-level artefacts, so a reader can substitute an independently authored label set without modifying the harness.

Table 11 maps each reported result to the stage that emits it and the recorded artefact it is computed from. One command, `experiments/run_baselines.sh`, executes the scanners, writes the raw output, and regenerates every row; `paper/Makefile` then refuses to typeset if any generated table is absent, so a manuscript built from this package cannot contain a results table that no recorded measurement supports; figures quoted in the prose are computed from the same recorded output but are not checked mechanically. The mapping is given because describing a package as complete is not the same as telling a reader which file answers which question.

9 Conclusion

This paper contributes an evaluation methodology for IaC security validation tools: a finding-normalization scheme that makes heterogeneous scanner output commensurable, an explicit treatment of matching strictness, a stated and re-runnable corpus admissibility rule whose value as a filter is untested here, and a statistical protocol suited to paired small-sample comparison. The accompanying framework composes repository-edge validation, native module testing, and plan-level policy evaluation, and is released with a harness in which every reported figure is a generated artefact traceable to recorded tool output.

Table 11 Each reported result, the stage that emits it, and the recorded artefact it derives from.

Result	Emitted by	Computed from
Corpus admissibility (Sec. 4.1)	corpus.py	benchmark/, results/ corpus_report.json
Criterion sensitivity (RQ1)	analyze.py	results/raw/
Effectiveness and intervals (RQ2)	analyze.py	results/raw/
Paired tests and $d_{\min}$ (RQ2)	analyze.py	results/evaluation.json
Latency and suite size (RQ3)	analyze.py	results/run_manifest.json, infrastructure/tests/
Layer attribution (RQ4)	analyze.py	results/raw/
Labelled external cases (RQ5)	analyze.py	results/raw/
Unlabelled subset (RQ5)	external.py	results/raw/external/, results/external_subset. json
Both diagrams	generate_figures.py	results/evaluation.json, results/run_manifest.json

benchmark/	
benchmark.json	planned case catalogue
generate_corpus.py	case generator
internal/	authored cases
external/	CIS-derived cases, unlabelled subset manifest
labelling/	blind second pass
evaluation/	
corpus.py	admissibility gate
normalize.py	finding normalization
control_map.json	canonical taxonomy
stats.py	exact CI + McNemar
tables.py	LaTeX table emitters
analyze.py	analysis entry point
external.py	unlabelled subset
security_framework/	
policies/*.rego	plan-level policy set
engine/	layer 1 scanning
infrastructure/	
tests/*.tftest.hcl	native module suite
results/	
raw/	per-case tool output
run_manifest.json	resolved versions
tables/	generated tables
pre_opa_polarity_fix/	pre-correction run
experiments/	
run_baselines.sh	full re-measurement
generate_figures.py	figure sources
fetch_external.py	pinned external repos

Fig. 3 Replication package layout. `pre_opa_polarity_fix` retains the measurements taken before the policy correction disclosed in Section 7.2.

The methodological findings are the transferable result. On this corpus the matching criterion moves reported recall by 29 to 43 percent of the spread between the tools compared, and by most of the available range for a tool whose findings are of a different kind from the corpus’s controls, which means comparative claims in this area should report their criterion as part of their method. Taxonomy coverage can favour the tool the taxonomy was authored around. Plan-level policy evaluation inherits a three-way serialisation distinction that silently inverts controls when it is collapsed. And corpus size is a weaker indicator of evidential strength than interval width and the smallest difference the chosen test could have resolved.

The outcome stated in Section 1 holds without adjustment: a third-party scanner attains the highest point-estimate control-level recall and the pipeline’s scored layers do not lead the comparison. That is one indication the procedure is not self-flattering, not proof of its neutrality, since Section 7.2 discloses a tuning asymmetry running the other way.

Three lines of future work follow directly. A mapping between the security-smell taxonomy that existing manually annotated IaC oracles are labelled against and the CIS control identifiers this evaluation uses would allow effectiveness to be measured on real-world configurations the investigator did not author, which is one of the mitigations Section 7.1 identifies as outstanding and the single change that would most strengthen the evidence. Enlarging the compliant baseline would narrow the specificity intervals to a width at which false-positive behaviour can be compared rather than merely reported. And extending the taxonomy beyond Terraform would establish whether the criterion sensitivity measured here is a property of IaC scanning or of this technology.

Statements and Declarations

Funding This research received no external funding.

Competing interests The author designed and implemented the plan-level policy set and the repository-edge layer evaluated here alongside third-party tools. This is a non-financial interest, and Sections 7.1 and 7.2 disclose the asymmetries it creates. The author declares no financial interests and is not affiliated with, and receives no support from, the vendor or maintainer of any tool evaluated in this study.

Ethics approval Not applicable. This study analyses software artefacts and involves no human participants or personal data.

Consent to participate Not applicable.

Consent for publication Not applicable.

Data and code availability The corpus, ground-truth annotations, canonical control taxonomy, normalization engine, execution harness, statistical implementation, analysis code, and the raw per-case tool output from which every table in this paper is generated are openly available:

<https://github.com/mchittineni/IaCSecBench>
<https://doi.org/10.5281/zenodo.21645016>

Author contributions The sole author designed the study, implemented the harness and the normalization engine, conducted the measurements, performed the analysis, and wrote the manuscript, with AI assistance as described under Use of AI-assisted technologies.

Use of AI-assisted technologies AI assistance was used in three roles. As a measuring instrument, a large language model was the second labeller in the automated consistency audit of Section 4.3, which states that the model’s identity, parameters and instruction were not recorded and what the resulting agreement does and does not establish in consequence; the per-case labels and reasons are in the replication package. In building the study, AI coding assistants were used to write parts of the harness and analysis code, the plan-level policy set and the benchmark cases, and to help debug the harness and inspect results, as described in Section 4.5. In preparing the manuscript, an AI assistant was used for language editing and to draft revisions of manuscript passages under the author’s direction. The assistant products used were Claude Code (Anthropic) and Gemini (Google); their model versions and parameters were not recorded. The study design and conclusions are the author’s. The author reviewed all AI-assisted work and accepts full responsibility for the manuscript and its content.

A Scope of the Validated Threats

The controls the pipeline enforces were selected against a STRIDE decomposition (spoofing, tampering, repudiation, information disclosure, denial of service, elevation of privilege) (Shostack, 2014) of a continuous-integration deployment path, summarised in Table 12. Protected assets are infrastructure configurations and compiled plan states, ingested datasets, policy and test definitions, and pipeline credentials; the trust boundaries separate the developer workspace from version control, the evaluation runtime from the target repository, and the compiled plan from the policy engine.

We record the decomposition to delimit which misconfiguration classes the corpus exercises, and claim nothing further from it. The pipeline was not subjected to adversarial testing and no adversary model was validated, so the table states intended coverage rather than a result; it is not a threat model of IaC pipelines, which this paper does not attempt.

Four domains are out of scope: post-deployment runtime compromise, including container and hypervisor vulnerabilities; supply-chain compromise of provider binaries or registry mirrors; IaC technologies other than Terraform, since the plan-level layer presupposes Terraform plan JSON; and physical and social-engineering attacks.

Table 12 STRIDE categories, mitigations, and the layer that enforces them.

Category	Mitigation	Layer
Spoofing	Authenticated API authorisers; transport encryption enforcement; prohibition of embedded credentials	1, 3
Tampering	Schema validation; pre-commit enforcement; evaluation of compiled plan documents	1, 2
Repudiation	Structured pipeline telemetry; recorded test execution; compliance tagging	2, 3
Information disclosure	Pattern-based personal-data detection; mandatory encryption at rest; storage public-access blocking	1, 3
Denial of service	Validation of web application firewall association, header handling, and request throttling	2, 3
Elevation of privilege	Policy denial of wildcard identity actions and unrestricted ingress ranges	3

References

- Agresti A, Coull BA (1998) Approximate is better than “exact” for interval estimation of binomial proportions. *The American Statistician* 52(2):119–126, <https://doi.org/10.1080/00031305.1998.10480550>
- Arcuri A, Briand L (2014) A hitchhiker’s guide to statistical tests for assessing randomized algorithms in software engineering. *Software Testing, Verification and Reliability* 24(3):219–250, <https://doi.org/10.1002/stvr.1486>
- Armbrust M, Fox A, Griffith R, Joseph AD, Katz R, Konwinski A, Lee G, Patterson D, Rabkin A, Stoica I, Zaharia M (2010) A view of cloud computing. *Communications of the ACM* 53(4):50–58, <https://doi.org/10.1145/1721654.1721672>
- Bass L, Weber I, Zhu L (2015) *DevOps: A Software Architect’s Perspective*. Addison-Wesley, Boston, MA, USA
- Böhme M, Szekeres L, Metzman J (2022) On the reliability of coverage-based fuzzer benchmarking. In: *Proc. 44th Int. Conf. Software Engineering (ICSE)*, pp 1621–1633, <https://doi.org/10.1145/3510003.3510230>
- Center for Internet Security (2021) CIS Amazon Web Services foundations benchmark, v1.4.0. https://www.cisecurity.org/benchmark/amazon_web_services, superseded; cited for requirement 2.1.1, removed in v3.0.0
- Center for Internet Security (2024) CIS Amazon Web Services foundations benchmark, v3.0.0. https://www.cisecurity.org/benchmark/amazon_web_services, accessed Aug. 2, 2026
- Chiari M, De Pascalis M, Pradella M (2022) Static analysis of infrastructure as code: A survey. In: *Proc. IEEE 19th Int. Conf. Software Architecture Companion (ICSA-C)*, pp 218–225, <https://doi.org/10.1109/ICSA-C54293.2022.00049>
- Clopper CJ, Pearson ES (1934) The use of confidence or fiducial limits illustrated in the case of the binomial. *Biometrika* 26(4):404–413, <https://doi.org/10.1093/biomet/26.4.404>
- Dalla Palma S, Di Nucci D, Tamburri DA (2020) Toward a catalog of software quality metrics for infrastructure code. *Journal of Systems and Software* 170:110726, <https://doi.org/10.1016/j.jss.2020.110726>
- Dalla Palma S, Di Nucci D, Palomba F, Tamburri DA (2022) Within-project defect prediction of infrastructure-as-code using product and process metrics. *IEEE Trans Software Engineering* 48(6):2086–2104, <https://doi.org/10.1109/TSE.2021.3051492>
- Dietterich TG (1998) Approximate statistical tests for comparing supervised classification learning algorithms. *Neural Computation* 10(7):1895–1923, <https://doi.org/10.1162/089976698300017197>
- Esposito M, Falaschi V, Falessi D (2024) An extensive comparison of static application security testing tools. In: *Proc. 28th Int. Conf. Evaluation and Assessment in Software*

- Engineering (EASE), pp 69–78, <https://doi.org/10.1145/3661167.3661199>
- Foaleam PL, Silva LD, Khomh F, Merlo E (2026) An empirical study of policy-as-code adoption in open-source software projects. *Journal of Systems and Software* 242:113028, <https://doi.org/10.1016/j.jss.2026.113028>
- Ge X, Fang C, Li X, Sun W, Wu D, Zhai J, Lin SW, Zhao Z, Liu Y, Chen Z (2023) Machine learning for actionable warning identification: A comprehensive survey. *arXiv preprint arXiv:2312.00324*, <https://doi.org/10.48550/arXiv.2312.00324>
- Gokhale AB, Jayaprakash V, Singaravelu N (2026) Comparative analysis of infrastructure-as-code misconfiguration detection through policy driven evaluation and taint-aware static reasoning. In: *Proc. Int. Conf. Computing, Communication and Networking Technologies (Lecture Notes in Networks and Systems)*, pp 400–413, https://doi.org/10.1007/978-3-032-27160-0_37
- Guerriero M, Garriga M, Tamburri DA, Palomba F (2019) Adoption, support, and challenges of infrastructure-as-code: Insights from industry. In: *Proc. IEEE Int. Conf. Software Maintenance and Evolution (ICSME)*, pp 580–589, <https://doi.org/10.1109/ICSME.2019.00092>
- Habib A, Pradel M (2018) How many of all bugs do we find? a study of static bug detectors. In: *Proc. 33rd IEEE/ACM Int. Conf. Automated Software Engineering (ASE)*, pp 317–328, <https://doi.org/10.1145/3238147.3238213>
- HashiCorp (2024) Terraform tests. <https://developer.hashicorp.com/terraform/language/tests>, accessed Aug. 2, 2026
- Hasselbring W (2021) Benchmarking as empirical standard in software engineering research. In: *Proc. 25th Int. Conf. Evaluation and Assessment in Software Engineering (EASE)*, pp 365–372, <https://doi.org/10.1145/3463274.3463361>
- Holm S (1979) A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics* 6(2):65–70, <https://www.jstor.org/stable/4615733>
- Humble J, Farley D (2010) *Continuous Delivery: Reliable Software Releases through Build, Test, and Deployment Automation*. Addison-Wesley, Boston, MA, USA
- Kapetanidou IA, Nizamis A, Votis K (2025) An evaluation of commonly used Kubernetes security scanning tools. In: *Proc. 2nd Int. Workshop on MetaOS for the Cloud-Edge-IoT Continuum (MECC)*, companion of EuroSys, pp 20–25, <https://doi.org/10.1145/3721889.3721924>
- Kumara I, Garriga M, Romeu AU, Di Nucci D, Palomba F, Tamburri DA, van den Heuvel WJ (2021) The do’s and don’ts of infrastructure code: A systematic gray literature review. *Information and Software Technology* 137:106593, <https://doi.org/10.1016/j.infsof.2021.106593>
- Li K, Chen S, Fan L, Feng R, Liu H, Liu C, Liu Y, Chen Y (2023) Comparison and evaluation on static application security testing (sast) tools for Java. In: *Proc. 31st ACM Joint European Software Engineering Conf. and Symp. Foundations of Software Engineering (ESEC/FSE)*, pp 921–933, <https://doi.org/10.1145/3611643.3616262>
- Matthews BW (1975) Comparison of the predicted and observed secondary structure of T4 phage lysozyme. *Biochimica et Biophysica Acta (BBA) – Protein Structure* 405(2):442–451, [https://doi.org/10.1016/0005-2795\(75\)90109-9](https://doi.org/10.1016/0005-2795(75)90109-9)
- McNemar Q (1947) Note on the sampling error of the difference between correlated proportions or percentages. *Psychometrika* 12(2):153–157, <https://doi.org/10.1007/BF02295996>
- Minna F, Blaise A, Tuma K, Massacci F (2026) Automated analysis of security policy violations in Helm charts. *IEEE Trans Dependable and Secure Computing* 23(2):2519–2533, <https://doi.org/10.1109/TDSC.2025.3628213>
- Morris K (2020) *Infrastructure as Code: Dynamic Systems for the Cloud Age*, 2nd edn. O’Reilly Media, Sebastopol, CA, USA
- National Institute of Standards and Technology (2020) Security and privacy controls for information systems and organizations. *Tech. Rep. Special Publication 800-53, Rev. 5*, NIST, <https://doi.org/10.6028/NIST.SP.800-53r5>
- Opdebeeck R, Zerouali A, De Roover C (2022) Smelly variables in Ansible infrastructure code: Detection, prevalence, and lifetime. In: *Proc. 19th Int. Conf. Mining Software Repositories (MSR)*, pp 61–72, <https://doi.org/10.1145/3524842.3527964>

- Opdebeeck R, Zerouali A, De Roover C (2023) Control and data flow in security smell detection for infrastructure as code: Is it worth the effort? In: Proc. 20th IEEE/ACM Int. Conf. Mining Software Repositories (MSR), pp 534–545, <https://doi.org/10.1109/MSR59073.2023.00079>
- Open Policy Agent (2024) Open Policy Agent documentation. <https://www.openpolicyagent.org/docs>, accessed Aug. 2, 2026
- OWASP Foundation (2023) OWASP benchmark project. <https://owasp.org/www-project-benchmark/>, accessed Aug. 2, 2026
- Qiu Y, Kon PTJ, Beckett R, Chen A (2024) Unearthing semantic checks for cloud infrastructure-as-code programs. In: Proc. 30th ACM Symp. Operating Systems Principles (SOSP), pp 574–589, <https://doi.org/10.1145/3694715.3695974>
- Rahman A, Mahdavi-Hezaveh R, Williams L (2019a) A systematic mapping study of infrastructure as code research. *Information and Software Technology* 108:65–77, <https://doi.org/10.1016/j.infsof.2018.12.004>
- Rahman A, Parnin C, Williams L (2019b) The seven sins: Security smells in infrastructure as code scripts. In: Proc. 41st Int. Conf. Software Engineering (ICSE), Montreal, QC, Canada, pp 164–175, <https://doi.org/10.1109/ICSE.2019.00033>
- Rahman A, Farhana E, Parnin C, Williams L (2020) Gang of eight: A defect taxonomy for infrastructure as code scripts. In: Proc. 42nd Int. Conf. Software Engineering (ICSE), pp 752–764, <https://doi.org/10.1145/3377811.3380409>
- Rahman A, Rahman MR, Parnin C, Williams L (2021) Security smells in Ansible and Chef scripts: A replication study. *ACM Trans Software Engineering and Methodology* 30(1):1–31, <https://doi.org/10.1145/3408897>
- Rahman A, Shamim SI, Bose DB, Pandita R (2023) Security misconfigurations in open source Kubernetes manifests: An empirical study. *ACM Trans Software Engineering and Methodology* 32(4):1–36, <https://doi.org/10.1145/3579639>
- Ralph P (2021) ACM SIGSOFT empirical standards released. *ACM SIGSOFT Software Engineering Notes* 46(1):19, <https://doi.org/10.1145/3437479.3437483>
- Roe H, Gogate M, Dashtipour K (2026) Multicloud security assessment: A benchmark study of infrastructure as code scanners. *ICCK Trans Information Security and Cryptography* 2(2):109–118, <https://doi.org/10.62762/tisc.2026.777114>
- Saavedra N, Ferreira JF (2022) GLITCH: Automated polyglot security smell detection in infrastructure as code. In: Proc. 37th IEEE/ACM Int. Conf. Automated Software Engineering (ASE), pp 1–12, <https://doi.org/10.1145/3551349.3556945>
- Saavedra N, Gonçalves J, Henriques M, Ferreira JF, Mendes A (2023) Polyglot code smell detection for infrastructure as code with GLITCH. In: Proc. 38th IEEE/ACM Int. Conf. Automated Software Engineering (ASE), pp 2042–2045, <https://doi.org/10.1109/ASE56229.2023.00162>
- Shamim MSI, Bhuiyan FA, Rahman A (2020) XI commandments of Kubernetes security: A systematization of knowledge related to Kubernetes security practices. In: Proc. IEEE Secure Development Conf. (SecDev), pp 58–64, <https://doi.org/10.1109/SecDev45635.2020.00025>
- Sharma T, Fragkoulis M, Spinellis D (2016) Does your configuration code smell? In: Proc. 13th Int. Conf. Mining Software Repositories (MSR), pp 189–200, <https://doi.org/10.1145/2901739.2901761>
- Shostack A (2014) *Threat Modeling: Designing for Security*. John Wiley & Sons, Indianapolis, IN, USA
- Tahaei M, Vaniea K, Beznosov KK, Wolters MK (2021) Security notifications in static analysis tools: Developers’ attitudes, comprehension, and ability to act on them. In: Proc. ACM CHI Conf. Human Factors in Computing Systems, pp 1–17, <https://doi.org/10.1145/3411764.3445616>
- Vassallo C, Panichella S, Palomba F, Proksch S, Gall HC, Zaidman A (2019) How developers engage with static analysis tools in different contexts. *Empirical Software Engineering* 25(2):1419–1457, <https://doi.org/10.1007/s10664-019-09750-5>
- Verdet A, Hamdaqa M, Da Silva L, Khomh F (2023) Exploring security practices in infrastructure as code: An empirical study. *arXiv preprint arXiv:2308.03952*, <https://doi.org/10.48550/arXiv.2308.03952>

- Verdet A, Hamdaqa M, Da Silva L, Khomh F (2025) Assessing the adoption of security policies by developers in Terraform across different cloud providers. *Empirical Software Engineering* 30(2), <https://doi.org/10.1007/s10664-024-10610-0>
- War A, Nikiema SL, Samhi J, Klein J, Bissyandé TF (2025) Security smells in infrastructure as code: A taxonomy update beyond the seven sins. *arXiv preprint arXiv:2509.18761*, <https://doi.org/10.48550/arXiv.2509.18761>
- Wohlin C, Runeson P, Höst M, Ohlsson MC, Regnell B, Wesslén A (2012) *Experimentation in Software Engineering*. Springer, Berlin, Germany, <https://doi.org/10.1007/978-3-642-29044-2>